

Midterm Report-Phase2

1. *ER-to-Relational (4 pt)*: Transform the ER diagram into relational model using SQL data definition language (DDL)

Turn boxes into tables (with simple primary keys).

Every entity in the ER diagram became its own table: customer, employee, supplier, product, and “order”. Each of these have their own primary key according to the ER diagram that I used for the primary key in the creation of the tables.

Map attributes straight to columns (and split composites).

Regular attributes just become columns. The only “special” case was the Customer Address, which is composite in ER-land. In a relational table you can’t keep a nested structure, so I flattened it into street, city, state, and zip to keep the first normal form. I also added basic constraints that match the data’s reality: non-negative stock and prices, valid order statuses, and a default order_date = NOW() so inserts don’t have to set a timestamp every time. I made customer.email and supplier.supplier_name unique to cut down on duplicates.

Convert one-to-many relationships into foreign keys.

The “many” side carries the pointer to the “one” side:

- product.supplier_id → supplier.supplier_id
- “order”.customer_id → customer.customer_id
- “order”.employee_id → employee.employee_id (nullable, since an order can exist even if we don’t track the handler)
I chose referential actions that match how the system should behave: deleting a supplier is restricted if products still reference it; deleting an employee sets null on existing orders (history stays intact); key updates cascade so FKs don’t break.

Replace many-to-many with a join table.

Order ↔ Product is M:N in the ER model, so I introduced order_line(order_id, product_id, quantity, line_price) with a composite primary key (order_id, product_id). That prevents duplicate lines for the same product in the same order. The FK to “order” is ON DELETE CASCADE (lines go away with the order), and the FK to product is RESTRICT (we don’t

want to delete a product that appears in old orders). Capturing quantity and line_price at the line level also “freezes” the transaction price.

Keep it normalized and usable.

Overall the schema lands in 3NF: each table is about one thing, attributes depend on the key, and repeating groups live in order_line. The keys, FKs, checks, and defaults carry over the ER diagram’s structure and rules, but in a way that’s practical for loading data and writing queries in PostgreSQL.

```

DROP SCHEMA IF EXISTS shop CASCADE;
CREATE SCHEMA shop;
SET search_path = shop;

CREATE TABLE customer (
    customer_id      INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    first_name       VARCHAR(50) NOT NULL,
    last_name        VARCHAR(50) NOT NULL,
    email            VARCHAR(255) UNIQUE,
    phone_number     VARCHAR(25),
    -- composite Address attribute flattened
    street           VARCHAR(120),
    city             VARCHAR(80),
    state            VARCHAR(50),
    zip              VARCHAR(15)
);

CREATE TABLE employee (
    employee_id      INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    first_name       VARCHAR(50) NOT NULL,
    last_name        VARCHAR(50) NOT NULL,
    position         VARCHAR(80),
    hire_date        DATE NOT NULL
);

CREATE TABLE supplier (
    supplier_id      INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    supplier_name    VARCHAR(120) NOT NULL,
    phone_number     VARCHAR(25),
    CONSTRAINT uq_supplier_name UNIQUE (supplier_name)
);

CREATE TABLE product (
    product_id       INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    product_name     VARCHAR(120) NOT NULL,
    category         VARCHAR(80),
    stock_quantity   INT NOT NULL DEFAULT 0 CHECK (stock_quantity >= 0),
    cost             NUMERIC(12,2) NOT NULL CHECK (cost >= 0),
    supplier_id      INT NOT NULL,
    CONSTRAINT fk_product_supplier
        FOREIGN KEY (supplier_id)
        REFERENCES supplier(supplier_id)
        ON UPDATE CASCADE
        ON DELETE RESTRICT
);

CREATE TABLE "order" (
    order_id          INT GENERATED ALWAYS AS IDENTITY PRIMARY KEY,
    order_date        TIMESTAMP NOT NULL DEFAULT NOW(),
    status            TEXT NOT NULL CHECK (status IN ('NEW','PAID','SHIPPED','CANCELLED')) DEFAULT 'NEW',
    payment_method    TEXT CHECK (payment_method IN ('CASH','CARD','TRANSFER','OTHER')),
    total_amount      NUMERIC(12,2) NOT NULL DEFAULT 0 CHECK (total_amount >= 0),
    customer_id       INT NOT NULL,
    employee_id       INT,
    CONSTRAINT fk_order_customer
        FOREIGN KEY (customer_id)
        REFERENCES customer(customer_id)
        ON UPDATE CASCADE
        ON DELETE RESTRICT,
    CONSTRAINT fk_order_employee
        FOREIGN KEY (employee_id)
        REFERENCES employee(employee_id)
        ON UPDATE CASCADE
        ON DELETE SET NULL
);

CREATE TABLE order_line (
    order_id          INT NOT NULL,
    product_id        INT NOT NULL,
    quantity          INT NOT NULL CHECK (quantity > 0),
    line_price        NUMERIC(12,2) NOT NULL CHECK (line_price >= 0),
    -- Composite PK = (order_id, product_id)
    PRIMARY KEY (order_id, product_id),
    CONSTRAINT fk_order_line
        FOREIGN KEY (order_id)
        REFERENCES "order"(order_id)
        ON UPDATE CASCADE
        ON DELETE CASCADE,
    CONSTRAINT fk_order_product
        FOREIGN KEY (product_id)
        REFERENCES product(product_id)
        ON UPDATE CASCADE
        ON DELETE RESTRICT
);

```

2. *Fill in the database with data* (synthetically generated, scraped from web data or integrated from open data) **(4 pt)**: . That should be done through a sequence of SQL Insert/Copy statements.

Data Source

All data were **synthetically generated** to simulate a realistic small-business scenario:

- suppliers.csv - list of product suppliers
- products.csv - inventory items linked to suppliers
- customers.csv - customer profiles with addresses
- employees.csv - store staff and hire dates
- orders.csv - customer orders referencing both customers and employees
- order_lines.csv - line-items associating products with orders

Import Method

We populated the schema via a sequence of SQL COPY statements executed in psql.

Example command sequence:

```
SET search_path TO shop;
```

```
COPY shop.supplier (supplier_name, phone_number)
FROM '/Users/abdullahalghabban/Desktop/SQLMidtermReport/suppliers.csv'
WITH (FORMAT csv, HEADER true);
```

```
COPY shop.product (product_name, category, stock_quantity, cost, supplier_id)
FROM '/Users/abdullahalghabban/Desktop/SQLMidtermReport/products.csv'
WITH (FORMAT csv, HEADER true);
```

```
COPY shop.customer (first_name, last_name, email, phone_number, street, city, state,
zip)
FROM '/Users/abdullahalghabban/Desktop/SQLMidtermReport/customers.csv'
WITH (FORMAT csv, HEADER true);
```

```
COPY shop.employee (first_name, last_name, position, hire_date)
```

```

FROM '/Users/abdullahalghabban/Desktop/SQLMidtermReport/employees.csv'
WITH (FORMAT csv, HEADER true);

COPY shop."order"  (order_date, status, payment_method, total_amount, customer_id,
employee_id)
FROM '/Users/abdullahalghabban/Desktop/SQLMidtermReport/orders.csv'
WITH (FORMAT csv, HEADER true);

COPY shop.order_line (order_id, product_id, quantity, line_price)
FROM '/Users/abdullahalghabban/Desktop/SQLMidtermReport/order_lines.csv'
WITH (FORMAT csv, HEADER true);

```

Verification

```

SELECT 'supplier',count(*) FROM shop.supplier
UNION ALL SELECT 'product',count(*) FROM shop.product
UNION ALL SELECT 'customer',count(*) FROM shop.customer
UNION ALL SELECT 'employee',count(*) FROM shop.employee
UNION ALL SELECT 'order',count(*) FROM shop."order"
UNION ALL SELECT 'order_line',count(*) FROM shop.order_line;

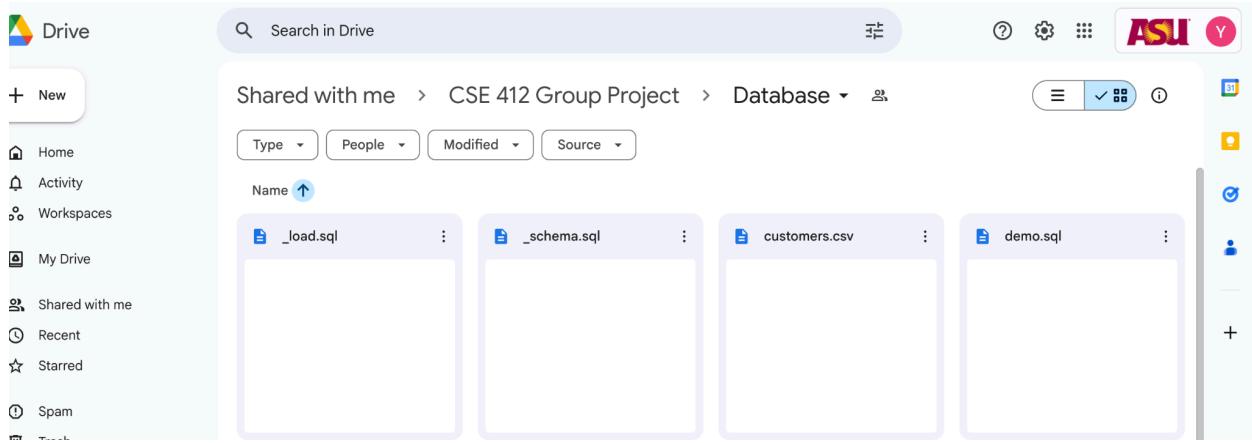
```

```

abdullahalghabban=# SELECT 'supplier',count(*) FROM shop.supplier
UNION ALL SELECT 'product',count(*) FROM shop.product
UNION ALL SELECT 'customer',count(*) FROM shop.customer
UNION ALL SELECT 'employee',count(*) FROM shop.employee
UNION ALL SELECT 'order',count(*) FROM shop."order"
[UNION ALL SELECT 'order_line',count(*) FROM shop.order_line;
 ?column? | count
-----+-----
 supplier |     8
 product  |    20
 customer |    10
 employee |     6
 order    |     8
 order_line|    27
(6 rows)

```

3.3. SQL Queries (4 pt): Prepare examples of SQL queries that cover the application description (as described by the proposal submitted). The SQL queries may include SELECT queries or INSERT/UPDATE/DELETE queries.



HOW TO RUN:

```
psql -d "$USER" -v ON_ERROR_STOP=1 -f _schema.sql
psql -d "$USER" -v ON_ERROR_STOP=1 -f _load.sql
psql -d "$USER" -v ON_ERROR_STOP=1 -f demo.sql
```

Code And Explain:

```
7   SELECT
8     (SELECT customer_id FROM customer ORDER BY customer_id LIMIT 1) AS cust_id,
9     (SELECT employee_id FROM employee ORDER BY employee_id LIMIT 1) AS emp_id,
10    (SELECT product_id FROM product ORDER BY product_id LIMIT 1) AS pid1,
11    (SELECT cost FROM product ORDER BY product_id LIMIT 1) AS price1,
12    (SELECT product_id FROM product ORDER BY product_id OFFSET 1 LIMIT 1) AS pid2,
13    (SELECT cost FROM product ORDER BY product_id OFFSET 1 LIMIT 1) AS price2;
14 \gset
15
16 \echo cust=:cust_id emp=:emp_id pid1=:pid1 pid2=:pid2 price1=:price1 price2=:price2
17
18 |
```

Sets the schema, then runs one SELECT that grabs a sample customer_id, employee_id, and two products with their current cost. \gset saves those columns into psql session variables (:cust_id, :emp_id, :pid1, :price1, :pid2, :price2). \echo prints the variables to confirm they're set.

```
SET
cust_id | emp_id | pid1 | price1 | pid2 | price2
-----+-----+-----+-----+
 1 | 1 | 1 | 79.00 | 2 | 129.00
(1 row)
```

cust=1 emp=1 pid1=1 pid2=2 price1=79.00 price2=129.00

The query returned one row and \gset created the variables. \echo shows: cust=1, emp=1, pid1=1, pid2=2, price1=79.00, price2=129.00. This means the variables are correctly initialized and ready to be used in later statements (e.g., INSERT ... VALUES (:cust_id, :emp_id, ...)).

```

19  -- SELECT #1: catalog by category
20  SELECT product_id, product_name, category, stock_quantity, cost
21  FROM product
22  WHERE category = 'Phones'
23  ORDER BY product_name;
24

```

product_id	product_name	category	stock_quantity	cost
1	Acme Phone X	Phones	228	79.00
2	Acme Phone XL	Phones	107	129.00
13	Acme Tablet 10	Phones	68	19.99
15	Stark Keyboard	Phones	64	1099.00

(4 rows)

Selects all products in the 'Phones' category and returns key fields—product_id, name, category, current stock_quantity, and cost—ordered by product_name. This implements the catalog “filter by category + show stock” behavior.

The output lists four phone items (IDs 1, 2, 13, 15) with their on-hand quantities and unit costs—for example, “Acme Phone X” (stock 228, cost 79.00) and “Acme Phone XL” (stock 107, cost 129.00). This confirms the filter works and displays real inventory and pricing for the category.

```

24
25  -- SELECT #2: product detail with supplier
26  SELECT p.product_id, p.product_name, p.category, p.stock_quantity, p.cost,
27  | | | s.supplier_name, s.phone_number
28  FROM product p
29  JOIN supplier s ON s.supplier_id = p.supplier_id
30  WHERE p.product_id = :pid1;
31

```

product_id	product_name	category	stock_quantity	cost	supplier_name	phone_number
1	Acme Phone X	Phones	228	79.00	Acme Electronics	555-7001

(1 row)

Joins product to its supplier to show a single product’s full detail (id, name, category, stock, cost) plus the supplier’s name and phone. The WHERE p.product_id = :pid1 filter uses the variable you set earlier, so it fetches the chosen product.

The result shows product 1 (“Acme Phone X”) in category “Phones,” with 228 units in stock at cost 79.00, supplied by “Acme Electronics” (phone 555-7001). This confirms the join works and the product-to-supplier relationship is correctly populated.

```
--  
32  -- INSERT: create a NEW order header  
33  INSERT INTO "order" (customer_id, employee_id, payment_method)  
34  VALUES (:cust_id, :emp_id, 'CARD')  
35  RETURNING order_id;  
36  \gset  
37
```

```
order_id
```

```
-----  
9
```

```
(1 row)
```

```
INSERT 0 1
```

```
INSERT 0 1
```

Creates a new order header using the previously set variables :cust_id and :emp_id, with payment_method = 'CARD'. The RETURNING order_id clause outputs the newly created ID, and \gset stores it into the psql variable :order_id for later statements.

The output shows one row returned with order_id = 9, and the INSERT 0 1 lines confirm exactly one row was inserted. Your session now has :order_id set to 9, ready for adding order lines and subsequent updates.

```
--  
38  -- INSERT: add two order lines (line_price uses current product cost)  
39  INSERT INTO order_line (order_id, product_id, quantity, line_price)  
40  SELECT :order_id, :pid1::INT, 2, :price1::NUMERIC  
41  UNION ALL  
42  SELECT :order_id, :pid2::INT, 1, :price2::NUMERIC;  
43
```

```
INSERT 0 2
```

Inserts two order-line rows for the current :order_id: one for :pid1 with quantity 2 and unit price :price1, and one for :pid2 with quantity 1 and unit price :price2. The two SELECT statements are combined with UNION ALL, so both rows are inserted in a single command. The message INSERT 0 2 confirms that exactly two rows were inserted into order_line, meaning both line items were added successfully for this order.

```
45  
44  -- UPDATE: change quantity for first line  
45  UPDATE order_line  
46  SET quantity = 3  
47  WHERE order_id = :order_id AND product_id = :pid1;  
48
```

```
UPDATE 1
```

Updates the existing line item for the current :order_id and product :pid1, setting its quantity to 3. This simulates a cart edit where the user increases the quantity of the first product.

UPDATE 1 means exactly one row matched the filter and was modified, so the quantity change took effect for that line item.

```
49  -- recompute order total
50  UPDATE "order" o
51  SET total_amount = COALESCE(
52    |   SELECT SUM(ol.quantity * ol.line_price)
53    |   FROM order_line ol
54    |   WHERE ol.order_id = o.order_id
55    |   ), 0
56  WHERE o.order_id = :order_id;
57
```

UPDATE 1

Recalculates the order header's total_amount as the sum of quantity * line_price across all its order_line rows, using COALESCE(..., 0) to set it to zero if there are no lines. The filter WHERE o.order_id = :order_id ensures only the current order is updated.

UPDATE 1 confirms exactly one order row was updated, meaning the header total now reflects the latest line quantities and prices (including the change you just made to set the first line's quantity to 3).

```
58  -- SELECT: order header after total recompute
59  SELECT order_id, order_date, status, payment_method, total_amount
60  FROM "order"
61  WHERE order_id = :order_id;
62
```

order_id	order_date	status	payment_method	total_amount
10	2025-10-25 13:49:04.543866	NEW	CARD	366.00

(1 row)

Selects the single order header by :order_id to verify the recalculated fields—showing the timestamp, current status, payment_method, and the updated total_amount after the line edits.

It returns order 10, still in NEW status with payment_method = CARD, and total_amount = 366.00, which matches $3 \times 79.00 + 1 \times 129.00$. This confirms the recompute worked and the header now reflects the cart change.

```
63  -- SELECT: order lines with line totals
64  SELECT ol.order_id, ol.product_id, p.product_name, ol.quantity, ol.line_price,
65  ||| (ol.quantity * ol.line_price) AS line_total
66  FROM order_line ol
67  JOIN product p ON p.product_id = ol.product_id
68  WHERE ol.order_id = :order_id
69  ORDER BY p.product_name;
70
```

order_id	product_id	product_name	quantity	line_price	line_total
10	1	Acme Phone X	3	79.00	237.00
10	2	Acme Phone XL	1	129.00	129.00

(2 rows)

Lists all line items for the current order by joining order_line with product to show the product name, quantity, unit price, and a computed line_total = quantity * line_price, ordered alphabetically by product name.

The result shows two items for order 10: “Acme Phone X” (qty $3 \times 79.00 = 237.00$) and “Acme Phone XL” (qty $1 \times 129.00 = 129.00$). The numbers add up to the header total 366.00, confirming line totals are correct.

```
71  -- SELECT #3: stock check for this order
72  SELECT ol.product_id, p.product_name,
73  ||| p.stock_quantity AS in_stock, ol.quantity AS needed,
74  ||| (p.stock_quantity - ol.quantity) AS remaining
75  FROM order_line ol
76  JOIN product p ON p.product_id = ol.product_id
77  WHERE ol.order_id = :order_id;
78
```

product_id	product_name	in_stock	needed	remaining
1	Acme Phone X	228	3	225
2	Acme Phone XL	107	1	106

(2 rows)

Selects a per-line stock snapshot for the current order by joining order_line with product. For each product it reports current on-hand units (in_stock), the order quantity (needed), and a

computed remaining = in_stock - needed, so you can validate availability before payment or fulfillment.

Both items have sufficient inventory: “Acme Phone X” 228 in stock (need 3 → remaining 225) and “Acme Phone XL” 107 in stock (need 1 → remaining 106). Non-negative remaining confirms no stock shortages for this order.

```
78  
79  -- UPDATE: decrement stock for each line  
80  UPDATE product p  
81  SET stock_quantity = p.stock_quantity - ol.quantity  
82  FROM order_line ol  
83  WHERE ol.order_id = :order_id  
84  | AND p.product_id = ol.product_id;  
85
```

UPDATE 2

Updates inventory in one set-based statement by joining product to order_line for the current :order_id, subtracting each line’s quantity from the matching product’s stock_quantity. The p.product_id = ol.product_id predicate ensures only items in this order are affected.

UPDATE 2 means two product rows were updated—both items in the order. Based on the prior stock check, the counts should now be reduced from 228→225 for “Acme Phone X” and 107→106 for “Acme Phone XL,” matching the ordered quantities.

```
86  -- UPDATE: set order status NEW -> PAID  
87  UPDATE "order"  
88  SET status = 'PAID'  
89  WHERE order_id = :order_id AND status = 'NEW';  
90
```

UPDATE 1

Updates the order’s status from NEW to PAID for the current :order_id, enforcing that only orders still in the NEW state are moved to PAID (prevents re-paying an already-paid/cancelled order).

UPDATE 1 indicates one order row matched the condition and was changed — the order is now marked PAID, which, together with the earlier stock decrements, represents a completed checkout for that order.

```

90
91  -- SELECT: verify header is PAID and show updated product stocks
92  SELECT order_id, status, total_amount
93  FROM "order"
94  WHERE order_id = :order_id;
95
96  SELECT p.product_id, p.product_name, p.stock_quantity
97  FROM product p
98  WHERE p.product_id IN (:pid1, :pid2)
99  ORDER BY p.product_id;
100

```

order_id	status	total_amount
10	PAID	366.00

(1 row)

product_id	product_name	stock_quantity
1	Acme Phone X	225
2	Acme Phone XL	106

(2 rows)

The first query retrieves the order header (order_id, status, total_amount) for the current :order_id to confirm the order's state and amount; the second query selects product_id, product_name, and current stock_quantity for the two products in the order to verify inventory levels after checkout.

The output shows order 10 is now PAID with a total of 366.00, and the two products' stocks have been reduced to 225 and 106 respectively, confirming the checkout updated both the order status and the product inventory.

```

100
101  -- SELECT #4: "My Orders" for the chosen employee
102  SELECT order_id, order_date, status, total_amount
103  FROM "order"
104  WHERE employee_id = :emp_id
105  ORDER BY order_date DESC;
106

```

order_id	order_date	status	total_amount
10	2025-10-25 13:49:04.543866	PAID	366.00
9	2025-10-25 13:49:04.537665	NEW	0.00
3	2025-09-23 12:00:00	PAID	0.00
2	2025-09-22 12:00:00	CANCELLED	0.00

```
(4 rows)
```

Selects all orders assigned to the chosen employee (:emp_id), returning order id, date, status, and total amount, ordered newest first — useful for an employee’s “My Orders” view to see work queue and recent activity.

The result shows four orders for that employee (most recent first): order 10 (PAID, 366.00), order 9 (NEW, 0.00), and two older orders (one PAID, one CANCELLED). This confirms the employee has both completed and pending orders to process.

```
106
107    -- UPDATE: advance PAID -> SHIPPED
108    UPDATE "order" SET status = 'SHIPPED'
109    WHERE order_id = :order_id AND status = 'PAID';
110
UPDATE 1
```

Updates the order’s status from PAID to SHIPPED for the current :order_id, enforcing that only orders already marked PAID are advanced to the shipping stage.

UPDATE 1 means one order row matched and was changed — the order has been transitioned to SHIPPED and is now in the fulfillment/shipping stage.

```
-- 
111    -- SELECT: verify shipped status
112    SELECT order_id, status FROM "order" WHERE order_id = :order_id;
113
order_id | status
-----+-----
  10 | SHIPPED
(1 row)
```

Updates the order row to fetch and display its current status for the given :order_id, used to confirm the latest lifecycle state after a status transition.

The query returns one row showing order_id = 10 with status = SHIPPED, indicating the order successfully moved from PAID to the shipping stage.

```
--  
114 -- INSERT: temporary customer with unique email (safe to re-run)  
115 WITH ins AS (  
116     INSERT INTO customer (first_name, last_name, email, phone_number, street, city, state, zip)  
117     SELECT 'Temp','User',  
118           'temp.user.' || to_char(clock_timestamp(),'YYYYMMDDHH24MISSMS') || '@example.com',  
119           NULL,'1 Demo St','DemoCity','DC','00000'  
120     RETURNING customer_id  
121 )  
122 SELECT customer_id FROM ins;  
123 \gset  
--  
customer_id  
-----  
11  
(1 row)
```

Inserts a temporary customer with a timestamped, unique email (so repeated runs won't violate the email uniqueness constraint) using a WITH ... RETURNING CTE, and then selects the new customer_id; the following \gset stores that returned customer_id into the psql variable :customer_id for immediate use.

The output shows the new customer_id = 11 (one row returned) and confirms :customer_id is set — you can now run DELETE FROM customer WHERE customer_id = :customer_id; to remove this temporary record.

```
125 -- DELETE: remove the temporary customer  
126 DELETE FROM customer  
127 WHERE customer_id = :customer_id;  
128
```

```
DELETE 1
```

Deletes the temporary customer row whose customer_id matches the :customer_id psql variable, removing the demo record from the customer table.

The output DELETE 1 confirms exactly one row was removed, so the temporary customer (the id returned earlier) was successfully deleted.

```
129  -- SELECT: verify deletion returns zero rows
130  SELECT customer_id, first_name, last_name
131  FROM customer
132  WHERE customer_id = :customer_id;
133  |
```

customer_id	first_name	last_name
(0 rows)		

Selects the customer row matching the temporary :customer_id to confirm whether the deletion succeeded; this query returns any remaining record fields (customer_id, first_name, last_name) for that id.

The result shows (0 rows), meaning no record exists with that customer_id — the temporary customer was successfully removed.